

Safe Globe — Product Requirements Document

Oct 1, 2026 · @h

Overview

Safe Globe is an AI-native global screening platform: AI agents run ongoing KYC on the organisations, people and vessels a company works with, and show the risk on an interactive map. It also gives screened organisations a voice: they can publish clarifications about sanctions or allegations against them.

Problem. Companies work with many stakeholders and partners around the world. Screening usually happens once at onboarding, but risk changes over time: new sanctions, new laws, fraud cases, and politically exposed persons (PEPs). Most companies have no easy way to keep screening continuously.

Solution. Two phases:

1. **Phase 1 — Screen and alert.** Screen organisations (including companies), people and vessels worldwide, and receive alerts on applicable laws and sanctions by country.
2. **Phase 2 — Clarify and engage.** Organisations in sanctioned countries, or accused of wrongdoing, can post statements explaining their position and the steps they are taking.

Goals and non-goals

Goals

- Use AI agents to screen any counterparty (organisation, person or vessel) against fraud, PEP and sanctions data, and clear false positives with a written reason.
- Keep screening ongoing, with alerts when a counterparty's risk or its country's laws and sanctions change.
- Show country risk visually on a map, so users understand exposure at a glance.
- Let users apply their own internal guidelines and policies on top of the base checks.
- Give users a path to deeper diligence by sending a DDQ (due diligence questionnaire) to a counterparty.
- In Phase 2, let screened organisations respond publicly with context and remediation steps.

Non-goals (for now)

- Making final compliance decisions for users. Safe Globe surfaces risk; the user decides.
- Verifying or endorsing organisation statements in Phase 2. They are shown as the organisation's own claims.
- Screening categories beyond fraud, PEP and sanctions in the first release (e.g. adverse media, ESG).

Target users

User	Who they are	What they need from Safe Globe
Screening company (primary)	Compliance, procurement, legal or risk teams at companies with global partners	Screen counterparties, get ongoing alerts, apply internal policies, send DDQs
Screened organisation (Phase 2)	A company or organisation that appears in screening results, e.g. one based in a sanctioned country	See how they appear, and publish clarifications and remediation steps
Individual (screened person)	Directors, owners or PEPs linked to an organisation	Be accurately represented; route to the organisation's response where relevant

What Safe Globe screens

Safe Globe screens three entity types — organisations (including companies), people and vessels — and links them in one ownership graph. A sanctioned person can then surface through the organisation they own and the ships that organisation operates.

Entity	Examples	Identifiers used for matching	Key risk checks	On the map
Organisations	Companies (suppliers, customers, distributors, partners), NGOs, charities, associations, state-owned entities, government bodies	Registration number, jurisdiction, addresses, directors and controlling persons	Sanctions (including ownership by sanctioned parties), terrorism-financing designations, state ownership, PEP links, fraud, adverse media	Pins at registered and operating addresses
People	Directors, owners and ultimate beneficial owners, PEPs, key contacts	Name variants and transliterations, date of birth, nationality, roles held	Sanctions, PEP status, fraud, adverse media	Shown through the organisations they are linked to, not by personal location
Vessels	Cargo ships, tankers, chartered and owned fleets	IMO number, MMSI, name history, flag, registered owner, manager and operator	Sanctioned vessel or owner, flag changes, AIS dark periods, ship-to-ship transfers, port calls in sanctioned countries	Live position and voyage track

Vessels matter for any customer in shipping, commodities, logistics or trade finance, and Singapore's port makes this a strong local wedge. People are deliberately not tracked by location, to respect privacy laws.

Custom data sources

Users can plug their own data into screening alongside Safe Globe's built-in data, through three connector types: API, data feed or web. Every source gets a label, so each finding shows exactly where it came from.

Connector	How users add it	Typical uses	How it is read
API	Enter the endpoint URL and API key, then map the response fields (name, ID, country, list type)	A paid screening or sanctions data provider, an internal watchlist service, a KYC vendor	Queried live at each screen and on every re-screen
Data feed	Upload a file or schedule a feed (CSV, XML or JSON, by upload, SFTP or URL)	Internal blocklists, a regulator's list file, a customer's own list of restricted parties	Re-imported on a schedule; changes trigger re-screening of affected counterparties
Web	Add website URLs or pages to monitor	A local regulator's enforcement page, a country gazette, industry news sites	AI reads the pages on a schedule and extracts relevant entities and changes

Requirements for every source

- **Label.** Users name each source (e.g. "Group blocklist" or "Vendor X sanctions API"). The label appears on every finding, alert and decision record that uses it.
- **Entity coverage.** Users choose which entity types the source covers: organisations, people, vessels or all three.
- **Trust level.** Users set how much to trust the source. The confidence engine weights each source's findings by this setting.
- **Test connection.** A test call or sample import checks the source works and the field mapping is right before it goes live.
- **Health monitoring.** Users get an alert if a source fails, returns errors or goes stale, and affected findings show reduced confidence until it recovers.
- **Secure keys.** API keys are encrypted, never shown again after entry, and can be rotated or revoked; only admins can add or edit sources.
- **Conflicts shown.** Where a custom source disagrees with built-in data, both findings are shown with their labels, so a reviewer decides.

Safe Globe will also offer its own screening API, so customers can screen counterparties directly from their onboarding, procurement or ERP systems.

Baseline: review workflow

Before any AI differentiation, Safe Globe must cover the standard third-party review workflow end to end. These are table stakes for compliance and risk buyers.

#	Capability	Requirement
B.1	Counterparty intake	A business owner introduces a new counterparty (name, service provided, country of registration) and submits it for review
B.2	Due diligence collection	Configurable forms collect the counterparty's answers and supporting documents (company extract, certificates, screening reports) in one record
B.3	Follow-up and correspondence	Follow-up questions, counterparty replies and documents stay attached to the review, so nothing is chased over email
B.4	Screening findings	Screening results and potential matches sit inside the review and stay open until a person has reviewed the evidence
B.5	Roles and ownership	Each review has a named owner (analyst) and reviewer (approver), with clear handoffs between them
B.6	Assessment drafts	AI drafts an editable assessment from the review record; a person edits and approves it
B.7	Decision record	The decision, the reasons behind it and the evidence used are saved with the assessment
B.8	Records search	Users ask questions across review records, e.g. outstanding work, who owns it, and the evidence behind a past decision
B.9	Oversight and reporting	A dashboard shows every review's status, owner and outstanding work, with reports built from case data
B.10	Configurable workflows	Teams bring their own questionnaires, evidence requirements and review steps
B.11	Access controls and security	Role-based permissions for sensitive cases; target SOC 2 Type II, ISO 27001 and GDPR alignment

The AI layer below builds on this baseline: Smart DDQ upgrades B.3, Sentry upgrades B.4, Case writer upgrades B.6, and Ask the Globe upgrades B.8.

AI layer

Safe Globe runs a team of AI agents that screen, monitor, investigate and draft, so a compliance officer reviews conclusions instead of raw hits. Every agent shows its evidence and sources; a person makes the final call.

Agent	What it does	Why it is an edge	Phase
Sentry (screening agent)	Screens counterparties against sanctions, PEP, fraud and adverse media in any language. Resolves entities using registration number, address, directors and dates (for vessels: IMO number, name history and owners), then auto-clears obvious false positives with a written reason	Cuts the manual false-positive work that dominates screening	1
Radar (regulatory agent)	Reads sanctions lists, government gazettes and news per country, in local languages. Maps each change to affected counterparties and writes a plain-language "what this means for you" alert	Alerts arrive with impact already worked out, not just a list update	1
Web (ownership graph)	Builds ownership and control networks from company registries and filings. Flags hidden links, e.g. a supplier 50% owned by a sanctioned person through two shell companies	Finds risk that name screening alone misses; drawn as arcs on the map	1

Agent	What it does	Why it is an edge	Phase
Forecast (predictive risk)	Scores how likely a country or counterparty is to become higher risk, using signals such as news sentiment, sanctions momentum and corruption-index trends	Moves users from reacting to sanctions to preparing for them; shown as a "risk weather" layer on the map	1 (beta)
Policy compiler	Reads a user's uploaded guidelines and policies and turns them into rules checked against every counterparty. Flags gaps and conflicts between policies	Custom policies become automatic checks, not a PDF someone has to remember	1
Ask the Globe	Natural-language questions over the map and records, e.g. "show suppliers within two ownership hops of a Russian state-owned company"	Queries the world, not only past review records	1
Smart DDQ	Writes a DDQ tailored to the specific findings, reads the answers, cross-checks them against registries and documents, and asks follow-ups where answers are vague or inconsistent	Due diligence that verifies, not just collects	1
Doc forensics	Checks uploaded certificates and extracts for tampering, mismatched metadata or forged stamps	Catches fake documents before they enter the record	1
Case writer	Drafts the decision memo with every claim linked to its source, ready for a reviewer to edit and sign	Audit-ready record in minutes	1

Agent	What it does	Why it is an edge	Phase
Truth check (statement verifier)	Checks Phase 2 statements against evidence, e.g. a claim of "operating before sanctions" against registry and sanction dates. Labels each claim consistent, unverified or contradicted	Stops sanctioned entities using statements to whitewash; makes the Phase 2 voice trustworthy	2
Tide (vessel agent)	Watches vessel AIS signals for dark periods, possible spoofing, ship-to-ship transfers and port calls in sanctioned countries. Links each vessel to its owners, managers and operators in the ownership graph	Catches sanctions evasion at sea that list screening alone misses; vessels move live on the map	1


AI principles

- **Human decides.** Agents recommend; a named reviewer approves every clearance and decision.
- **Show the work.** Every conclusion links to its sources and the data it used, with a confidence level.
- **Full audit trail.** Every agent action is logged, so a regulator can replay how a decision was made.
- **Governed models.** Bias and accuracy are tested regularly, in line with frameworks such as MAS FEAT principles and the EU AI Act.

Confidence engine

Every finding in Safe Globe carries two separate ratings: **risk level** (how serious it would be) and **confidence level** (how sure we are it is true). Together they decide what happens next, so analysts spend time only where it matters.

Risk sets the stakes; confidence sets who acts



Routing by risk level and confidence level · 4 actions

Only low-risk, high-confidence findings are cleared automatically; anything high-risk always goes to a person.

Where confidence levels appear

- **Screening matches.** Each potential match shows how likely it is to be the same entity, and why: e.g. matches on registration number and director, differs on address. Vessels match on IMO number, past names and owners.
- **Alerts.** Radar alerts show how sure the AI is that a law or sanctions change applies to this counterparty.
- **Ownership links.** Each link in the ownership graph shows whether it comes from an official registry, a filing, or an inference from news.
- **Statements and DDQ answers.** Truth check rates each claim as consistent, unverified or contradicted, with a confidence level.
- **Counterparty profile.** An overall confidence score shows how complete the evidence is. Countries with weak public registries lower it, and the map can show this as a data coverage layer.

How it works

Requirement	Detail
Confidence bands	High, Medium and Low, each backed by a percentage and the factors that drove it
Reasons shown	Every rating lists what raised it and what lowered it, with links to the evidence
Configurable thresholds	Each customer sets the confidence needed to auto-clear or escalate, matched to their risk appetite and policies
Raise confidence	For any Low or Medium rating, the AI suggests the next step that would raise it, e.g. request a company extract or send a DDQ question
Re-scoring	Confidence updates automatically when new evidence arrives (documents, DDQ answers, registry changes)
Calibration	The system tracks how often reviewers overturn AI ratings per band; High confidence must prove accurate before auto-clearing is switched on
Recorded in decisions	The confidence level at the time of decision is saved in the decision record for audit

Phase 1: Screen and alert (KYC)

Phase 1 delivers base KYC on fraud, PEP and sanctions, ongoing alerts, and a map-based interface.

#	Feature	Requirement	Priority
1.1	Counterparty screening	Search and screen organisations (including companies), people and vessels worldwide against fraud, PEP and sanctions lists	Must have
1.2	Ongoing monitoring	Save counterparties to a watchlist and re-screen them continuously, not just at onboarding	Must have

#	Feature	Requirement	Priority
1.3	Country law and sanctions alerts	Alert users when sanctions or applicable laws change in a country where they have counterparties	Must have
1.4	Map interface	A Google Maps-style world view showing counterparties by location, with country risk overlays (e.g. corruption perception scores)	Must have
1.5	Custom guidelines and policies	Users upload their own internal guidelines or policies, which add to the base fraud, PEP and sanctions checks	Should have
1.6	Company profile	A page per organisation where users can read more about it: screening results, location, country context	Must have
1.7	DDQ request	If a user is not satisfied with a profile, they can send a due diligence questionnaire to the organisation	Should have
1.8	Custom data sources	Users connect their own labelled API, data feed or web sources, which feed into screening alongside built-in data (see Custom data sources)	Must have

Interface direction. The map is the centre of the product. It should feel bold, modern and visually striking, not like a typical compliance dashboard. Users zoom into regions and click pins to open company profiles.

Reference data. Country-level risk can draw on Transparency International's [Corruption Perceptions Index 2025](#).

Phase 2: Clarify and engage


Phase 2 lets flagged organisations tell their side, so screening companies see context rather than just a red flag.

Example. Company A works with Company B. Company B is based in Russia, a sanctioned country. Company B posts a statement: it was operating there before the sanctions took effect, and lists the steps it is taking in response. Company A sees this statement on Company B's profile alongside the screening results.

#	Feature	Requirement	Priority
2.1	Organisation claim	An organisation can claim its profile after verifying it is the real entity	Must have
2.2	Clarification statements	A claimed organisation can publish statements about sanctions, allegations or other flags	Must have
2.3	Visibility to screeners	Statements appear on the profile, clearly labelled as the organisation's own claims	Must have
2.4	DDQ responses	Organisations can receive and answer DDQs sent by screening companies	Should have
2.5	Statement updates	Organisations can update statements as their situation changes; screeners watching them get an alert	Could have

Key user flow

Every alert ends in a review, and doubts become a DDQ



Safe Globe user flow · screening company and screened organisation

Phase 1 runs along the top: screen, monitor, get alerted, review. In Phase 2 the organisation's statements and DDQ answers feed back into that same profile review.

Data, risks, metrics and open questions

Data sources needed

- Sanctions lists (e.g. UN, US, EU, UK and local lists such as Singapore's)
- PEP databases
- Fraud and enforcement records
- Country risk indices, such as the Corruption Perceptions Index
- Law and regulation change feeds by country
- Vessel data: IMO and MMSI registries, AIS position tracking, flag and ownership history

Risks

- **Data accuracy and false positives.** Wrongly flagging an organisation or person could cause harm and legal exposure.
- **Data licensing cost.** Commercial sanctions and PEP data is expensive; this shapes pricing.
- **Phase 2 misuse.** Sanctioned entities could use statements to whitewash. Statements must be clearly labelled as unverified claims.
- **Privacy.** Screening individuals requires compliance with data protection laws (e.g. Singapore's PDPA, GDPR).
- **AI errors.** A wrongly auto-cleared match is the worst failure. Start with AI clearing only high-confidence cases, sample-check them, and keep a human sign-off on every decision.
- **Explainability.** Regulators and auditors will ask why the AI concluded something; every output must cite its sources.
- **Vessel data gaps.** AIS signals can be switched off or spoofed. Treat dark periods as a risk signal, and lower confidence when positions cannot be verified.

Success metrics

- Number of counterparties under ongoing monitoring
- Time from a sanctions or law change to a user alert
- Share of alerts users mark as relevant (false-positive rate)
- DDQs sent and response rate
- Phase 2: organisations that claim profiles and post statements
- Share of screening hits auto-cleared by AI, and the error rate found in sample checks

- Analyst hours per counterparty review, compared with manual screening
- Confidence calibration: share of High-confidence ratings that reviewers uphold (target above 95% before auto-clearing is enabled)

Open questions

- ☐ Which data providers to license for sanctions, PEP and fraud data?
- ☐ Who is the first target customer segment, and in which market?
- ☐ How will an organisation prove it owns a profile in Phase 2?
- ☐ Pricing model: per seat, per monitored counterparty, or per screen?
- ☐ What default confidence thresholds should customers start with, and who at the customer can change them?

Sources

- [Transparency International — Corruption Perceptions Index 2025](#)